

\$0.155 per 1,000 narratives · 0.20 s median · \$30.22 for the whole run

### THE AUDIT DECIDES WHAT p MEANS

**Coded CRIS fields**  
agreement, 150,000 crashes

**Human gold set**  
accuracy, 2,416 labels

**Calibration audit**  
z = -8.2: rejected

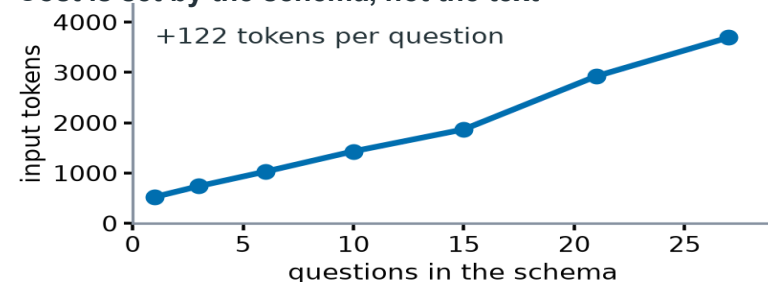
The coded fields make the model look better calibrated than it is, for every variable. Recalibrating on the human labels is what makes p usable.

**Baselines, held apart**  
same narratives, same criteria

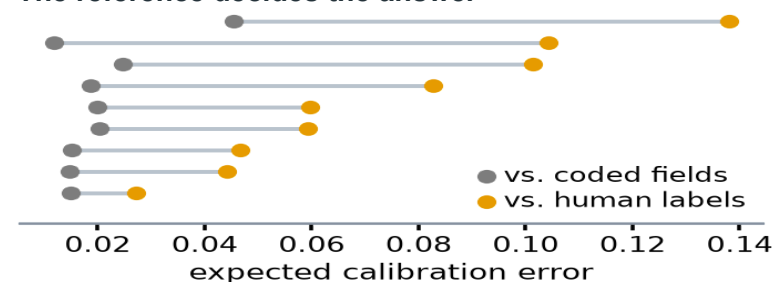
A frontier model gains +0.068  $F_1$  at 51× the price. A second gains nothing measurable. Neither is calibrated.

### WHAT IT FOUND

#### Cost is set by the schema, not the text



#### The reference decides the answer



#### The budget, variable by variable

